

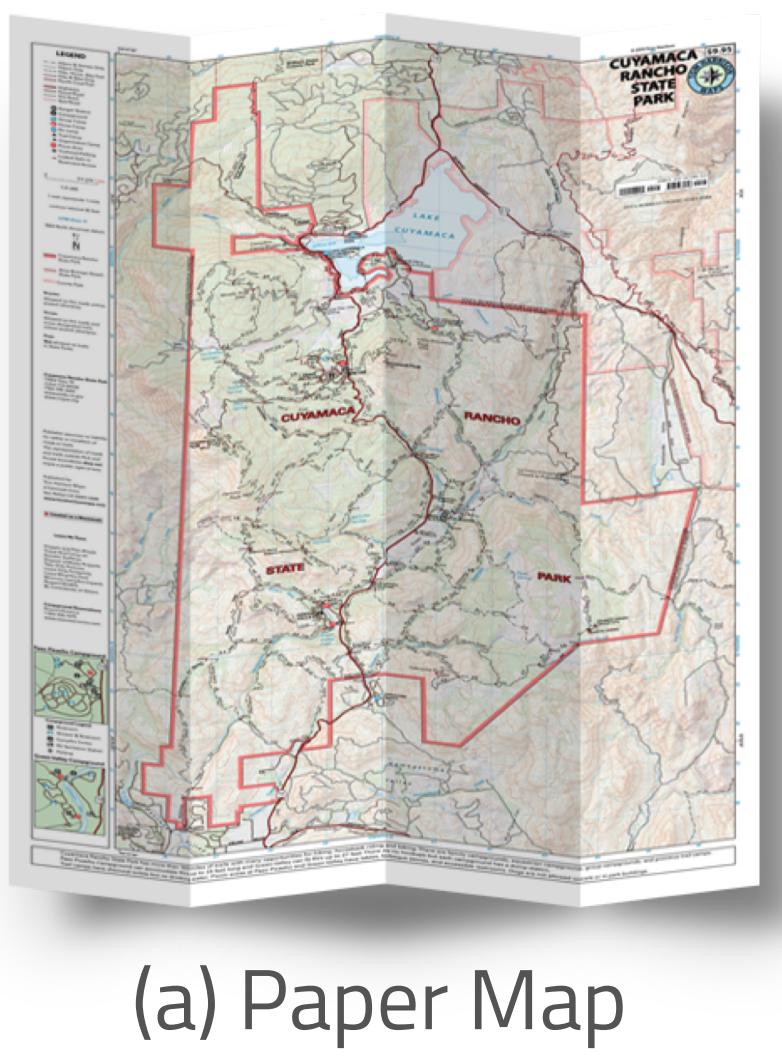


Distributed Processing Model for Temporal Graphs

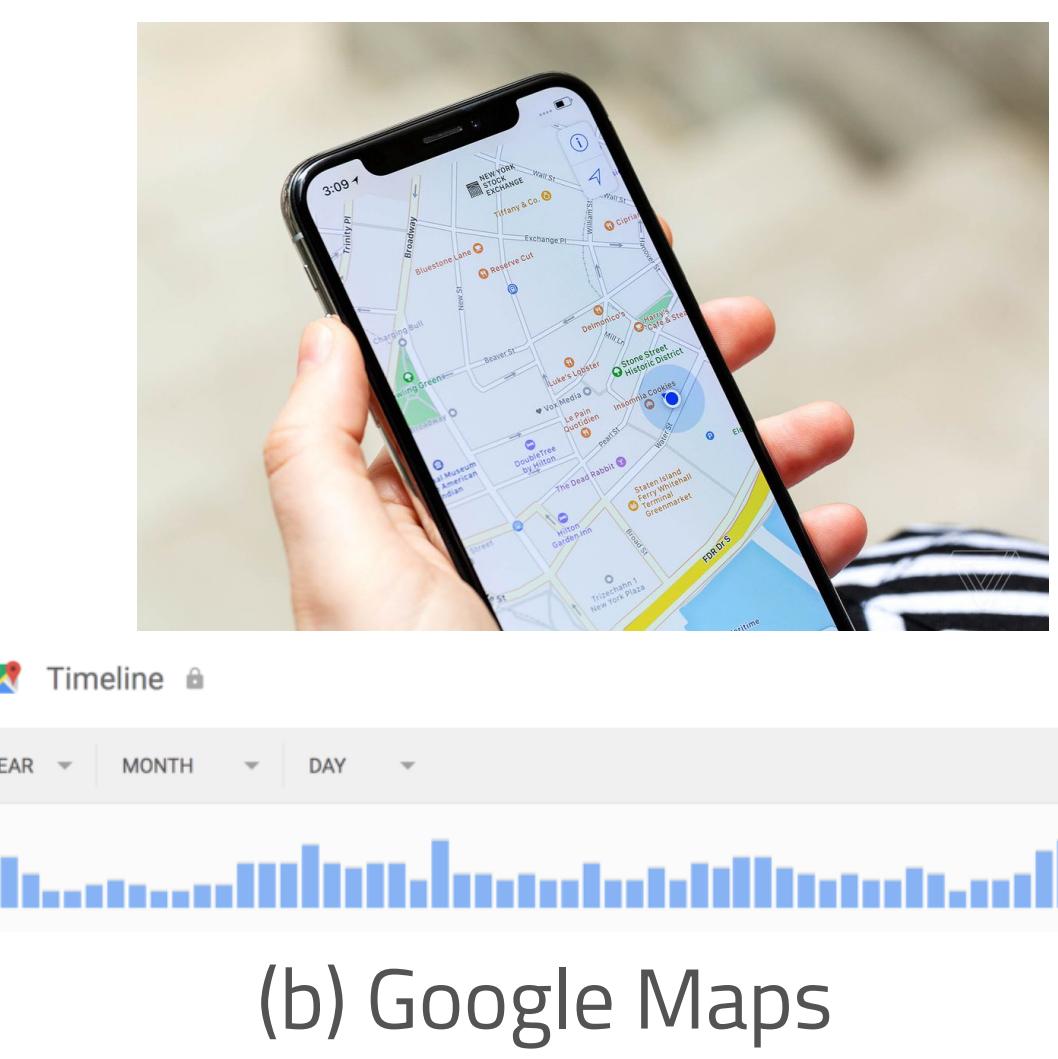
Swapnil Gandhi and Yogesh Simmhan

INTRODUCTION

- **Graphs** are widely considered to be the natural means of representation for many networks
- But real-world networks are often **evolving** with links being added/removed and properties updated over time



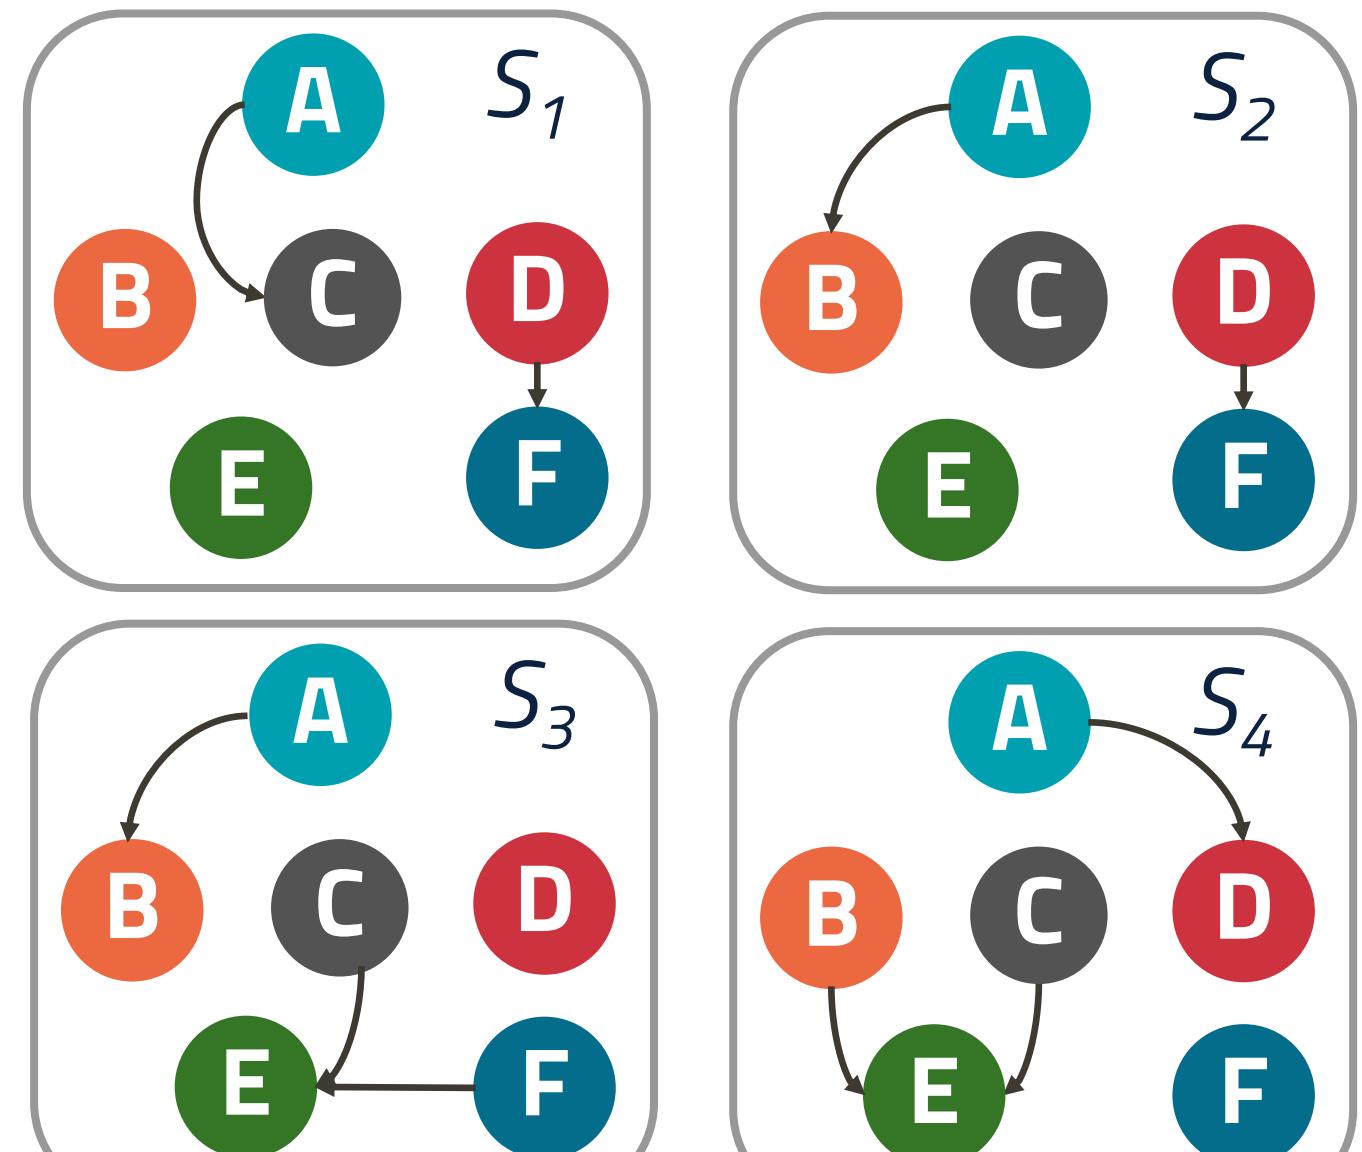
(a) Paper Map



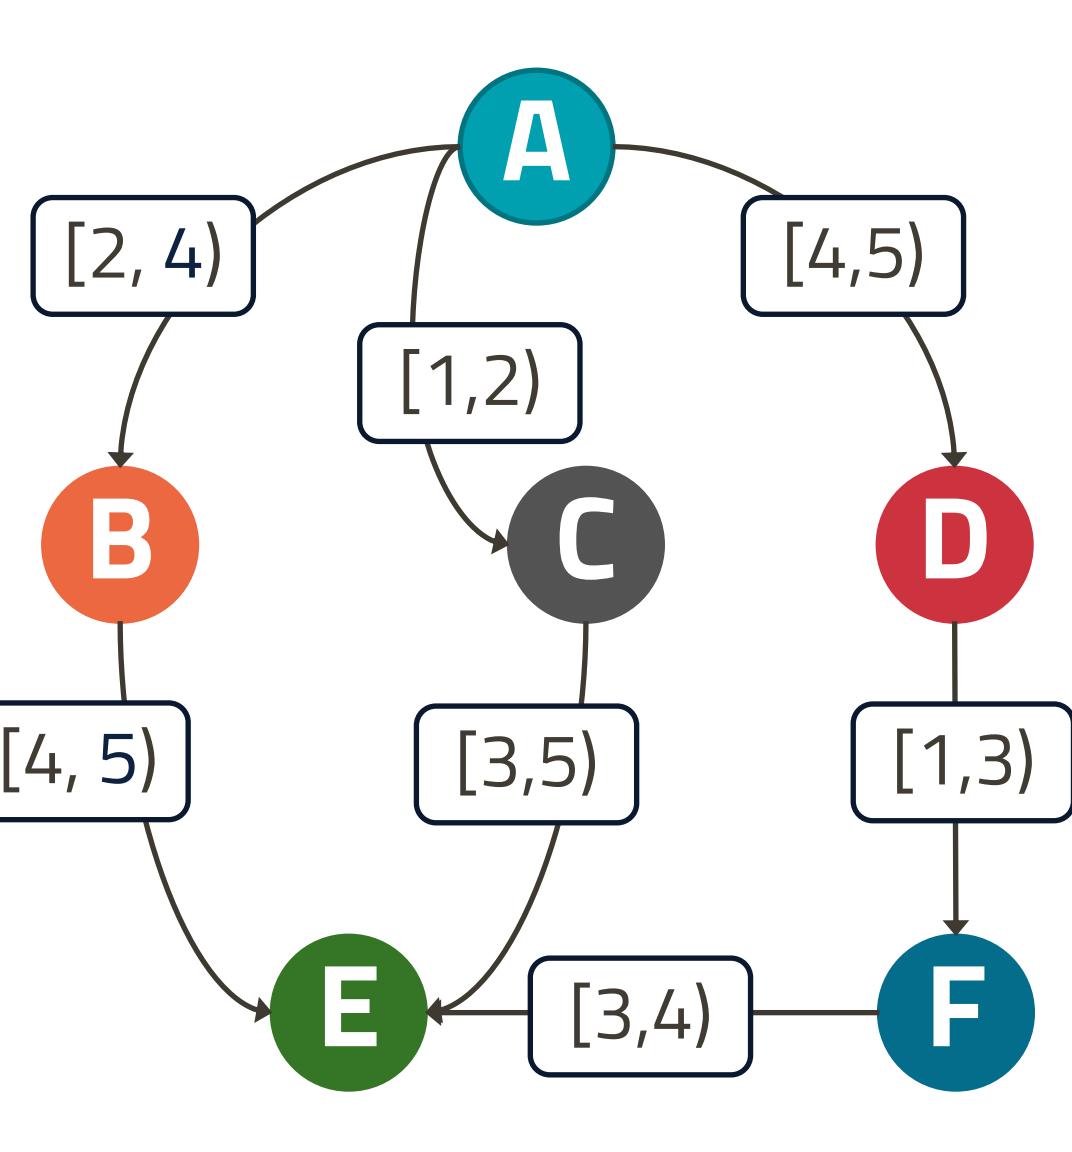
(b) Google Maps

- **Temporal Graphs** are dynamic networks which contain existence information for all vertices & edges at every point in time

Examples : Social, Citation/Collaboration, Sensor & Transit Network, Human Connectomes, Internet-of-Things ...



(c) Multi-Snapshot Graph



(d) Temporal Graph

CHALLENGES:

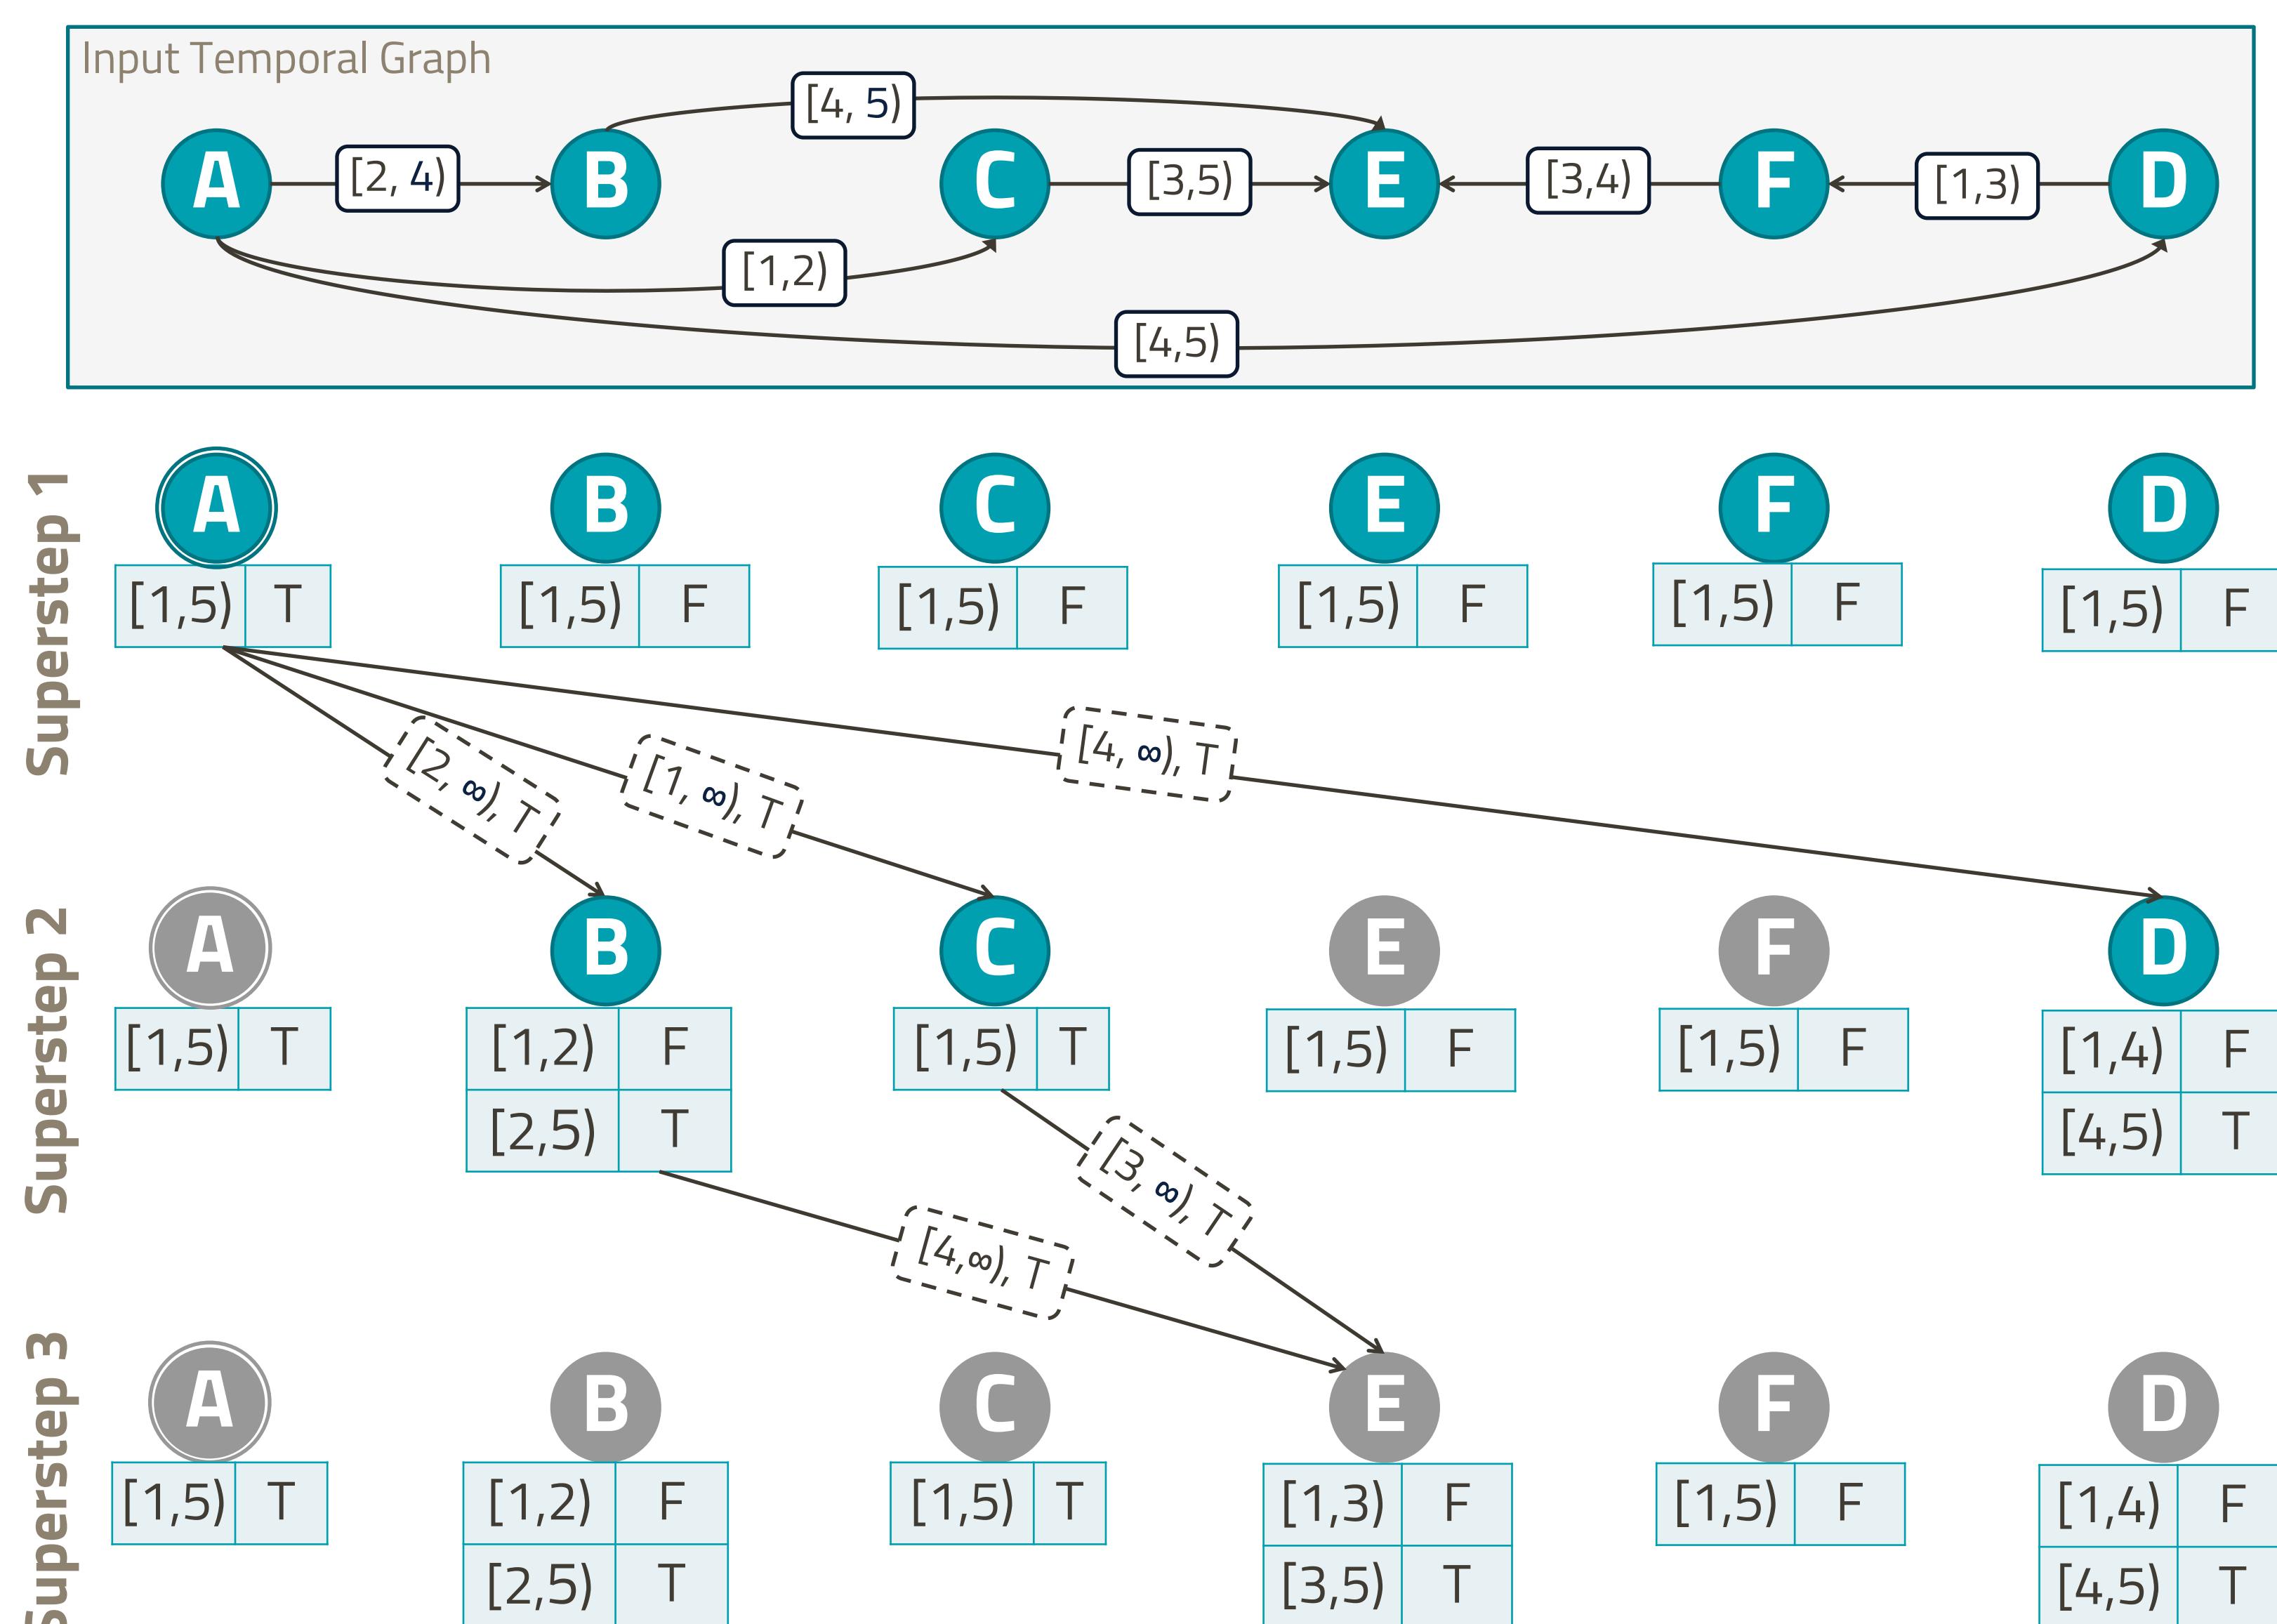
1. Lack of **unifying abstraction** to operate on temporal graphs
Existing abstractions either work only for time-dependent algorithms or time-agnostic algorithms but not for both classes

2. Existing abstractions either *do not scale* on distributed systems or are *not even designed* for **distributed execution**

Specialized native time-dependent algorithms are designed for single-threaded shared-memory execution

This Work : First **scalable, distributed and general-purpose** programming abstraction for processing arbitrarily large temporal graphs

TEMPORAL REACHABILITY



INTERVAL-CENTRIC COMPUTATION

- ICM extends vertex-centric computation model with Interval as a **first-class citizen**
 - ▶ Vertex (and edge) attributes, computation state and messages have **time-validity** and may have associated computation
- User writes program from the **perspective** of a single interval
 - ▶ Infrastructure lifts computation for arbitrary-sized graph

Compute method is executed at every active interval of a vertex in each superstep. Compute can inspect/modify state associated with active interval and has access to all interval messages received from previous superstep.

```
1 void init(Vertex v) {
2     v.setState(v.interval, false);
3 }
4
5 void compute(Vertex v, Interval t, boolean
vstate, Message[] msgs) {
6     if(getSuperstep() == 1 && isSource(v)) {
7         v.setState(t, true);
8     }
9     if(!msgs.isEmpty()) v.setState(t, true);
10 }
11
12 Message scatter(Edge e, Interval t, boolean
vstate) {
13     return new Message(e, new
Interval(t.start, infinity), true);
14 }
```

Initializes state of vertex v for its entire lifespan

Scatter permits user to send message with interval validity.
The function has read-only access to vertex state associated with interval and edge properties

TIME-WARP

- Allows user-logic to **consistently** operate over multiple vertex sub-intervals **in-parallel**
 - ▶ Reminiscent of **Temporal Aggregate**
 - ▶ Uses **one-pass** algorithm
 - ▶ Support for online aggregation

Messages	
State	
τ_m	M
$[0,4]$	m_1
$[0,5]$	s_1
$[2,7]$	m_2
$[5,9]$	s_2
$[5,7]$	m_3
$[9,10]$	s_3
$[5,9]$	m_4
$[9,10]$	m_5

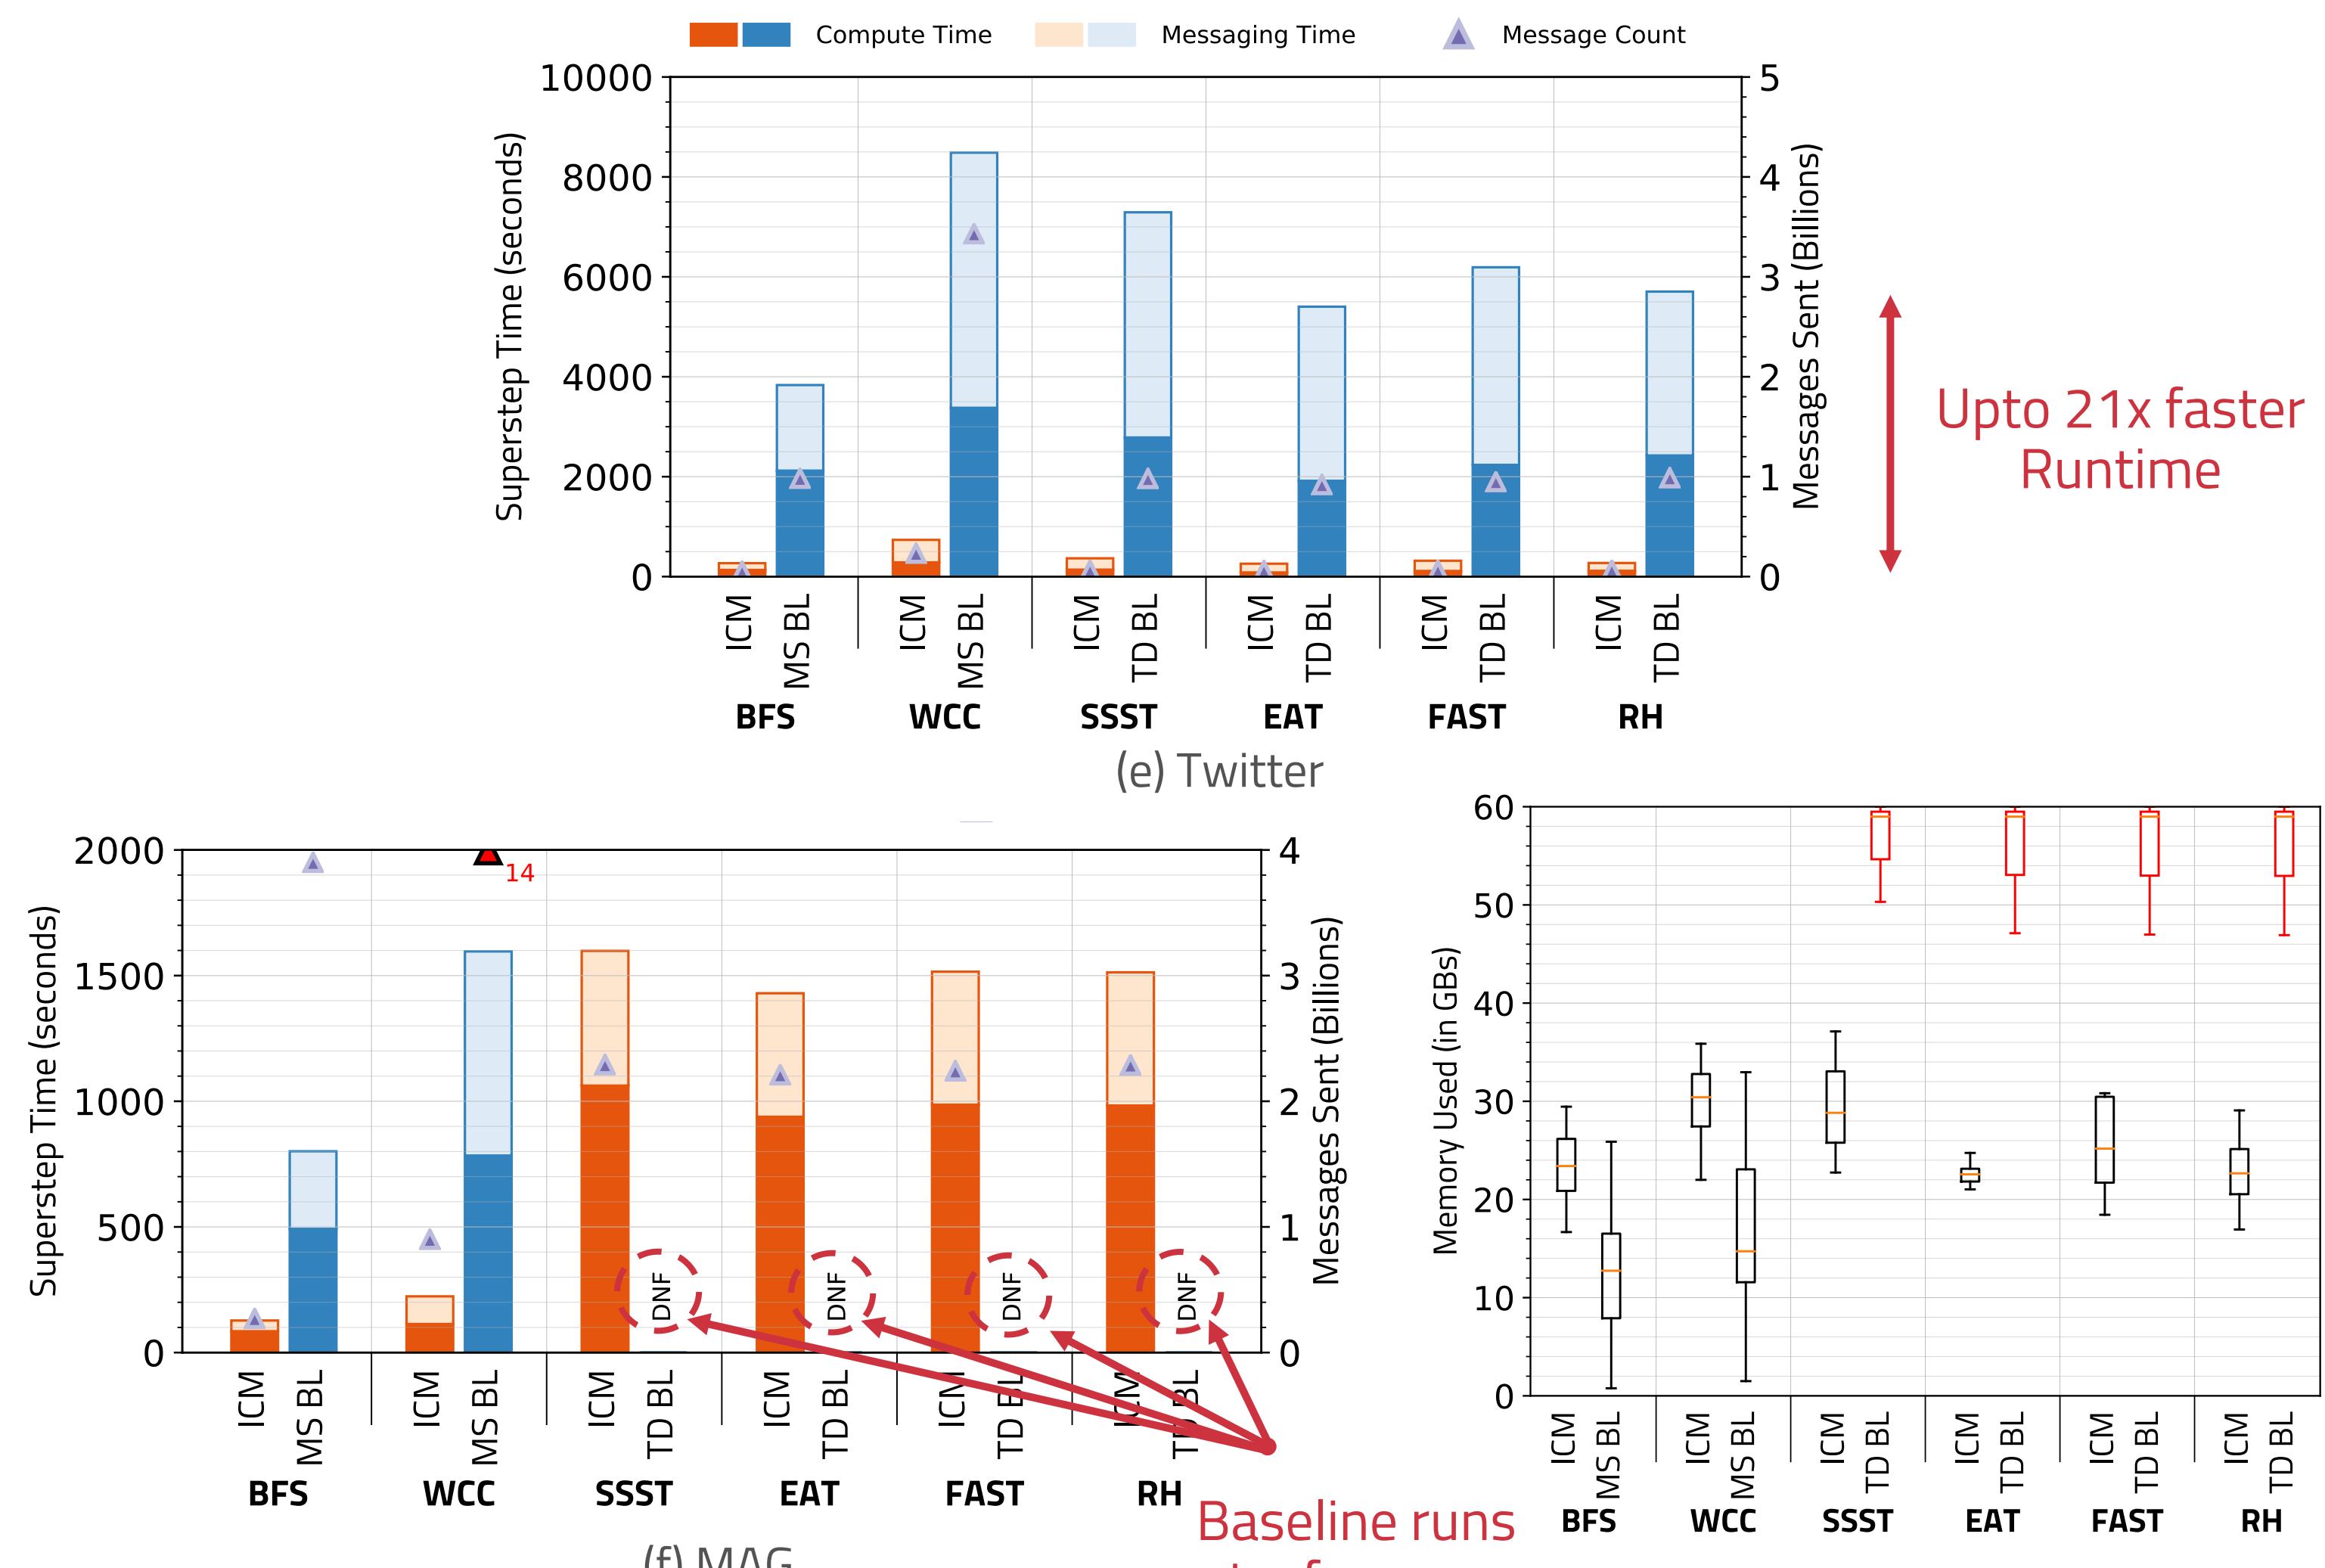
TW	S	M
w ₁	$[0,2]$	s_1
w ₂	$[2,4]$	s_1, m_2
w ₃	$[4,5]$	s_1
w ₄	$[5,7]$	s_2, m_2, m_3, m_4
w ₅	$[7,9]$	s_2
w ₆	$[9,10]$	s_3

Advantages :

- Transparantly performs temporal alignment, re-partitioning, replication and grouping
- **Minimizes** compute calls; **coalesces** interval-state and outgoing messages

EXPERIMENTAL EVALUATION

Dataset	Snapshots	Avg. Lifespan		ICM Graph		Multi-Snapshot Graph		Transformed Graph	
		V	E	V _{ICM}	E _{ICM}	Σ V	Σ E	V _{TD}	E _{TD}
Twitter	30	29.5	28.4	43.9M	2.01B	1.3B	60B	219M	8.3B
MAG	219	20.9	12.75	116M	1.02B	3.4B	13B	2.6B	9.6B



Longer Vertex and edge lifespan → Better performance

SUMMARY & ON-GOING WORK

- First **scalable API** for processing temporal graph
 - ▶ Leverages existing solutions & makes it practical
- On-going work towards adding support for real-time streaming updates and temporal graph partitioning

